

More data manipulation with dplyr and tidy

Completed by Leon Woltermann

Contents

Aims of this worksheet	1
Data joining with two table verbs (<code>left_join()</code> et al.)	2
Data reshaping (<code>spread()</code> and <code>gather()</code>)	9

NB: The worksheet has been developed and prepared by Lincoln Mullen. Source: Lincoln A. Mullen, *Computational Historical Thinking: With Applications in R (2018-)*: <http://dh-r.lincolnmullen.com>. Minor modifications added by Maxim Romanov (loading `methodists` dataset).

The best way to learn R or computational history is to practice. These worksheets contain a series of questions designed to teach you about R or different computational methods. The worksheets are R Markdown documents that include text and code together. The places where you are expected to answer questions are marked like this.

(©) Can you make a plot from this dataset?

Beneath each question is a space to either create a code block or write an answer.

Aims of this worksheet

In an earlier worksheet, you learned the basic data manipulation verbs from the dplyr package: `select()`, `filter()`, `mutate()`, `arrange()`, `group_by()`, and `summarize()`. In this worksheet you will learn additional data verbs from the dplyr and tidyr packages. These data verbs relate to window functions (`lead()` and `lag()`), data table joins (`left_join()` et al.), and data reshaping (`spread()` and `gather()`)

To begin, we will load the necessary packages, as well as the Methodist data.

```
library(tidyverse)
library(historydata)
#data("methodists")
#methodists
```

methodists data (MGR)

If `methodists` dataset does not load, we can try the following. First, restart R (in the menu: **Session > Restart R**), then run the following lines:

```
devtools::install_github("ropensci/historydata", force=TRUE)
```

```
## Downloading GitHub repo ropensci/historydata@HEAD
```

```
## * checking for file '/private/var/folders/6f/0x08zkks1754nb4kts9p0t240000gn/T/Rtmp0RxYtj/remotes2c8f
## * preparing 'historydata':
## * checking DESCRIPTION meta-information ... OK
## * checking for LF line-endings in source and make files and shell scripts
## * checking for empty or unneeded directories
## * building 'historydata_0.2.9001.tar.gz'
```

```
library(historydata)
data(methodists)
#methodists
```

Alternatively, we can load the data differently. The package itself is available on gitHub (<https://github.com/ropensci/historydata>), so we can try a different way of getting the data that we need for the worksheet. Specifically, if we know the exact address of the data file (url), we can open it with the `read.csv` command, like shown below (you need to be connected to Internet, of course):

```
methodists <- read.csv("https://raw.githubusercontent.com/ropensci/historydata/master/data-raw/methodists.csv")
#methodists
```

This data file, however, is slightly different from what we need, so some minor modifications will be necessary. You do not need to be concerned about the code in the next chunk, just run it.

```
#library(dplyr)

replace_na <- function(x, val = 0L) {
  ifelse(is.na(x), val, x)
}

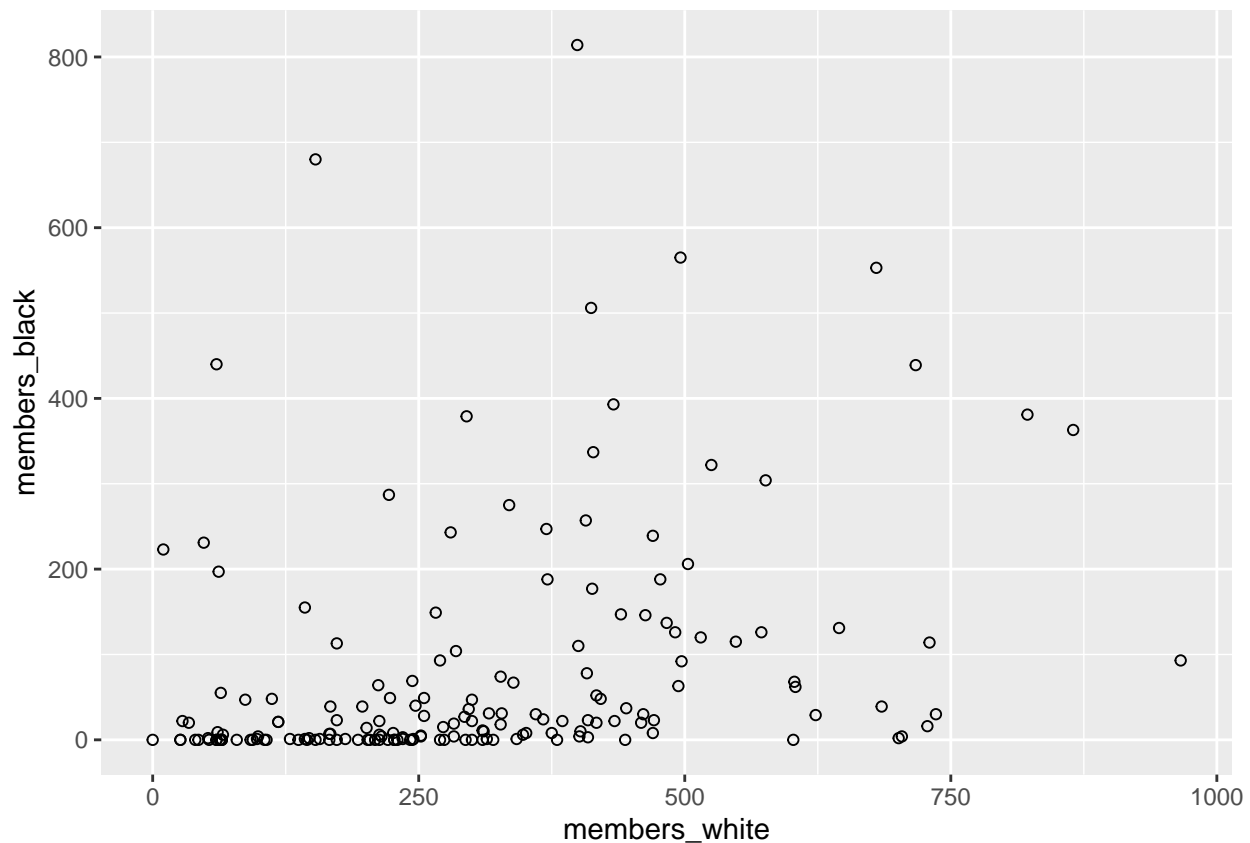
methodists <- methodists %>%
  as_tibble() %>%
  filter(minutes_year != 1778,
         minutes_year != 1779,
         minutes_year != 1785) %>%
  filter(minutes_year >= 1786, minutes_year <= 1834) %>%
  dplyr::rename(members_black = members_colored,
               year = minutes_year) %>%
  mutate(members_indian = as.integer(members_indian)) %>%
  mutate(members_white = replace_na(members_white),
         members_black = replace_na(members_black),
         members_indian = replace_na(members_indian)) %>%
  rowwise() %>%
  mutate(members_total = sum(members_general, members_white, members_black,
                             members_indian, na.rm = TRUE)) %>%
  ungroup() %>%
  select(year, conference, district, meeting, state, members_total,
         starts_with("members_"), url)
```

Data joining with two table verbs (`left_join()` et al.)

It is often the case that we want to use some variable in our data to create a new variable. Consider the Methodist data for the year 1800. Perhaps we are interested in the racial composition of the churches. Do they tend to be all white and all black, or do some churches have both white and black members in varying proportions? The simplest way to get a look at that question is to create a scatter plot of the figures for white and black membership.

```
methodists_1800 <- methodists %>%
  filter(year == 1800) %>%
  select(meeting, state, members_white, members_black)

ggplot(methodists_1800, aes(x = members_white, y = members_black)) +
  geom_point(shape = 1)
```



That scatterplot is interesting as far as it goes, but we might reasonably suspect that the racial composition of methodist meetings varies by region. We could use the `state` variable to facet the plot by state. However, this has two problems. There are 20 states represented in that year. Our faceted plot would have 20 panels, which is too many. But more important, by looking at individual states we might be getting *too* fine grained a look at the data. We have good reason to think that it is regions that matter more than states.

It is easy enough to describe what we would do to translate states into a new column with regions. We would look at each state name and assign it to a region. Connecticut would be in the Northeast, New York would be in the Mid-Atlantic, and so on. We can think of this problem as looking up a value in one table (our Methodist data) in another table. That other table will have a row for each state, where each state name is associated with a region. (In many cases, though, it would make more sense to create a CSV file with the data and read it in as a data frame.)

```
regions <- tibble(
  state = c("Connecticut", "Delaware", "Georgia", "Kentucky", "Maine",
            "Maryland", "Massachusetts", "Mississippi", "New Hampshire",
            "New Jersey", "New York", "North Carolina",
            "Northwestern Territory", "Pennsylvania", "Rhode Island",
            "South Carolina", "Tennessee", "Upper Canada", "Vermont",
            "Virginia"),
  region = c("Northeast", "Atlantic South", "Atlantic South", "West",
             "Northeast", "Atlantic South", "Northeast", "Deep South",
             "Northeast", "Mid-Atlantic", "Mid-Atlantic", "Atlantic South",
             "West", "Mid-Atlantic", "Northeast", "Atlantic South", "West",
             "Canada", "Northeast", "Atlantic South")
)
```

And now we can inspect the table.

```
regions
```

```
## # A tibble: 20 x 2
##   state      region
##   <chr>      <chr>
## 1 Connecticut Northeast
## 2 Delaware   Atlantic South
## 3 Georgia    Atlantic South
## 4 Kentucky   West
## 5 Maine      Northeast
## 6 Maryland   Atlantic South
## 7 Massachusetts Northeast
## 8 Mississippi Deep South
## 9 New Hampshire Northeast
## 10 New Jersey Mid-Atlantic
## 11 New York   Mid-Atlantic
## 12 North Carolina Atlantic South
## 13 Northwestern Territory West
## 14 Pennsylvania Mid-Atlantic
## 15 Rhode Island Northeast
## 16 South Carolina Atlantic South
## 17 Tennessee  West
## 18 Upper Canada Canada
## 19 Vermont    Northeast
## 20 Virginia   Atlantic South
```

We can do a look up where we take the `state` column in the `methodists_1800` data frame and associate it with the `states` column in our `regions` data frame. The result will be a new column `region`. Notice how we use the `by =` argument to specify which column in the left hand table matches which column in the right hand table.

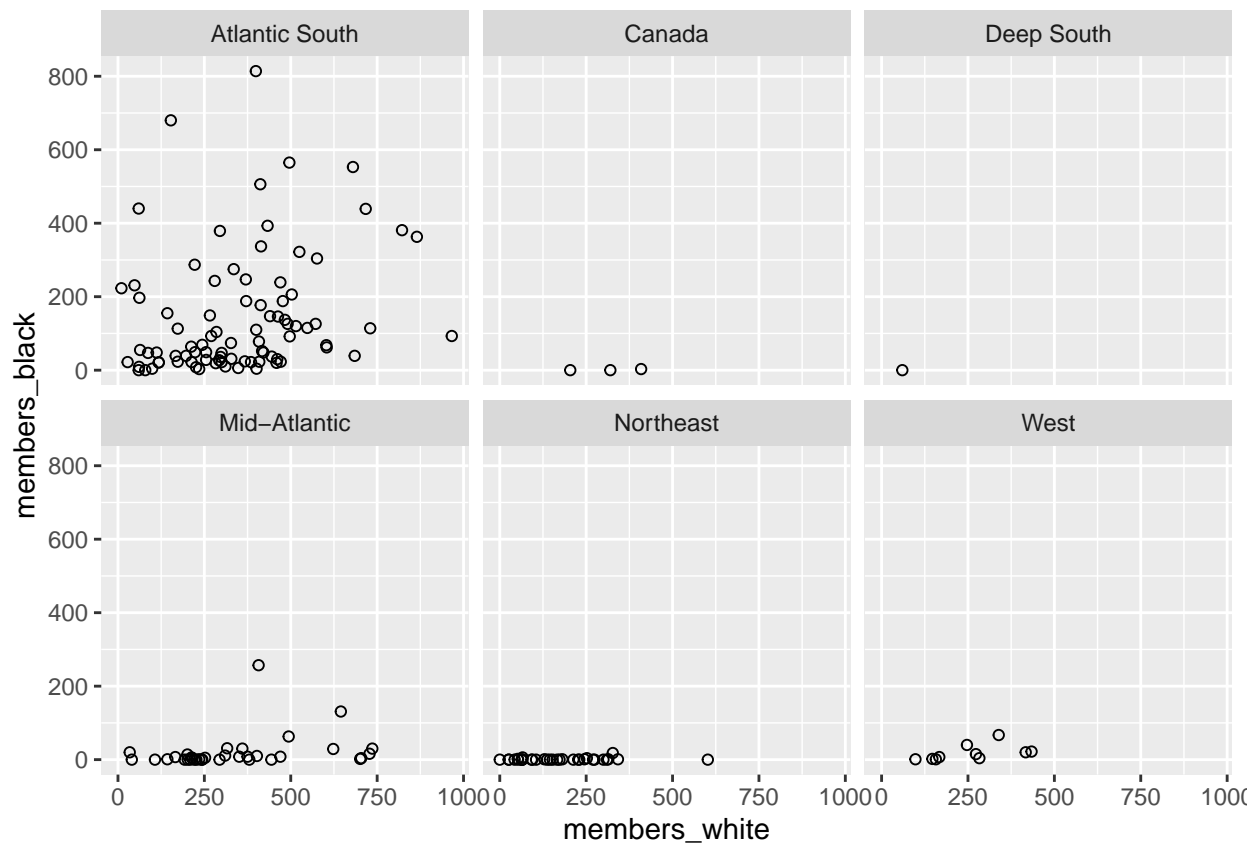
```
methodists_region <- methodists_1800 %>%
  left_join(regions, by = "state")
```

```
methodists_region
```

```
## # A tibble: 169 x 5
##   meeting      state      members_white members_black region
##   <chr>      <chr>          <int>          <int> <chr>
## 1 Augusta    Georgia             61             9 Atlantic South
## 2 Burke      Georgia            297            36 Atlantic South
## 3 Richmond   Georgia            548           115 Atlantic South
## 4 Washington Georgia            497            92 Atlantic South
## 5 Broad River South Carolina    604            62 Atlantic South
## 6 Bush River South Carolina    328            31 Atlantic South
## 7 Charleston South Carolina     60           440 Atlantic South
## 8 Cherokee   South Carolina     79             0 Atlantic South
## 9 Edisto      South Carolina    572           126 Atlantic South
## 10 Georgetown South Carolina     10           223 Atlantic South
## # ... with 159 more rows
```

Then we can plot the results. As we suspected, there is a huge regional variation.

```
ggplot(methodists_region, aes(x = members_white, y = members_black)) +
  geom_point(shape = 1) +
  facet_wrap(~ region)
```



- (1) Can you summarize the racial composition of the different regions by year (i.e., a region had a certain percentage white and black members for a given year) and create a plot of the changing racial composition in each region over time?

```
methodists_full <- methodists %>%
  select(year, state, members_total, members_white, members_black)

methodists_full_region <- methodists_full %>%
  left_join(regions, by="state") %>%
  select(-state)

methodists_full_region
```

```
## # A tibble: 20,241 x 5
##   year members_total members_white members_black region
##   <int>         <int>         <int>         <int> <chr>
## 1 1786             356             330             26 <NA>
## 2 1786             488             416             72 <NA>
## 3 1786             364             305             59 <NA>
## 4 1786             412             382             30 <NA>
## 5 1786             429             392             37 <NA>
## 6 1786             540             524             16 <NA>
## 7 1786             449             374             75 <NA>
## 8 1786             178             167             11 <NA>
## 9 1786             368             350             18 <NA>
## 10 1786            166             140             26 <NA>
## # ... with 20,231 more rows
```

```

racial_composition <- methodists_full_region %>%
  filter(!is.na(region)) %>%
  filter(members_total!=0) %>%
  group_by(year, region) %>%
  summarize(members_total = sum(members_total), members_white = sum(members_white), members_black = sum(members_black),
  mutate(percentage_white = members_white / members_total * 100, percentage_black = members_black / members_total * 100),
  mutate(relation = percentage_black / percentage_white * 100)

```

`summarise()` has grouped output by 'year'. You can override using the
`.groups` argument.

```
racial_composition
```

```


## # A tibble: 31 x 8
## # Groups:   year [6]
##   year region      members_total members_white members_black percentage_white
##   <int> <chr>          <int>          <int>          <int>          <dbl>
## 1 1796 Atlantic So~      42059          31521          10538          74.9
## 2 1796 Canada           474            472            2            99.6
## 3 1796 Mid-Atlantic      9406            8703            703            92.5
## 4 1796 Northeast        2519            2509            10            99.6
## 5 1796 West             2296            2169            127            94.5
## 6 1797 Atlantic So~     41976          30459          11517          72.6
## 7 1797 Mid-Atlantic     10513            9950            563            94.6
## 8 1797 Northeast        2999            2974            25            99.2
## 9 1797 West             2373            2274            99            95.8
## 10 1798 Atlantic So~    41822          30267          11555          72.4
## # ... with 21 more rows, and 2 more variables: percentage_black <dbl>,
## #   relation <dbl>

```

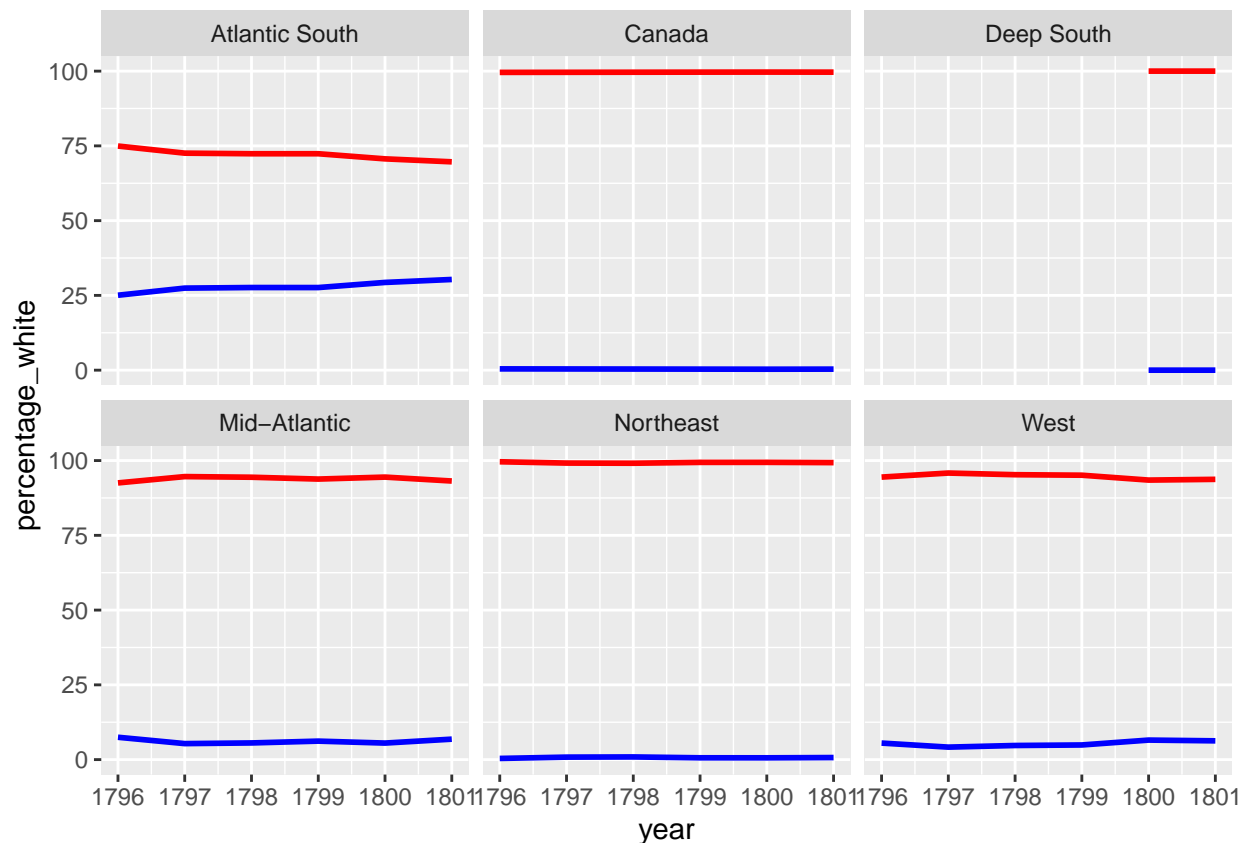
```

ggplot(racial_composition) +
  geom_line(aes(x = year, y = relation)) + facet_wrap(~ region)

```



```
ggplot(racial_composition) +
  geom_line(aes(x=year,y=percentage_white), size=1, col="red") +
  geom_line(aes(x=year,y=percentage_black), size=1, col="blue") +
  facet_wrap(~ region)
```



- (2) In the `europop` package there are two data frames, `europop` with the historical populations of European cities, and `city_coords` which has the latitudes and longitudes of those cities. Load that package and join the two tables together. Can you get the populations of cities north of 48° of latitude?

```
#devtools::install_github("mdlincoln/europop", force=TRUE)
```

```
library(europop)
```

```
data("europop")
```

```
merged_cities <- europop %>%
  left_join(city_coords, by="city") %>%
  filter(lat>=42)
```

```
merged_cities[merged_cities$lat > 42, ]
```

```
## # A tibble: 2,128 x 6
```

	city	region	year	population	lon	lat
	<chr>	<chr>	<int>	<int>	<dbl>	<dbl>
## 1	BERGEN	Scandinavia	1500	0	5.33	60.4
## 2	COPENHAGEN	Scandinavia	1500	NA	12.6	55.7
## 3	GOTEBORG	Scandinavia	1500	0	12.0	57.7
## 4	KARLSKRONA	Scandinavia	1500	0	15.6	56.2
## 5	OSLO	Scandinavia	1500	0	10.7	59.9
## 6	STOCKHOLM	Scandinavia	1500	0	18.1	59.3
## 7	BATH	England and Wales	1500	0	-2.36	51.4
## 8	BIRMINGHAM	England and Wales	1500	0	-1.90	52.5
## 9	BLACKBURN	England and Wales	1500	0	-2.48	53.8
## 10	BOLTON	England and Wales	1500	0	-2.43	53.6


```
## # ... with 2,118 more rows
```

- (3) In the `historydata` package there are two tables, `judges_people` and `judges_appointments`. Join them together. What are the names of black judges who were appointed to the Supreme Court?

```
judge_merge <- judges_people %>%
  left_join(judges_appointments, by="judge_id")

judge_merge[judge_merge$court_name == "Supreme Court of the United States" & judge_merge$race == "African American"]
```

```
## # A tibble: 2 x 27
##   judge_id name_first name_middle name_last name_suffix birth_date
##   <int> <chr>      <chr>      <chr>      <chr>      <int>
## 1    1489 Thurgood   <NA>      Marshall <NA>      1908
## 2    2362 Clarence   <NA>      Thomas   <NA>      1948
## # ... with 21 more variables: birthplace_city <chr>, birthplace_state <chr>,
## #   death_date <int>, death_city <chr>, death_state <chr>, gender <chr>,
## #   race <chr>, court_name <chr>, court_type <chr>, president_name <chr>,
## #   president_party <chr>, nomination_date <chr>, predecessor_last_name <chr>,
## #   predecessor_first_name <chr>, senate_confirmation_date <chr>,
## #   commission_date <chr>, chief_judge_begin <int>, chief_judge_end <int>,
## #   retirement_from_active_service <chr>, termination_date <chr>, ...
```

- (4) What courts did those justices serve on before the Supreme Court?

```
filter(judge_merge, judge_id == 1489 | judge_id == 2362, court_name != "Supreme Court of the United States")
```

```
## # A tibble: 2 x 27
##   judge_id name_first name_middle name_last name_suffix birth_date
##   <int> <chr>      <chr>      <chr>      <chr>      <int>
## 1    1489 Thurgood   <NA>      Marshall <NA>      1908
## 2    2362 Clarence   <NA>      Thomas   <NA>      1948
## # ... with 21 more variables: birthplace_city <chr>, birthplace_state <chr>,
## #   death_date <int>, death_city <chr>, death_state <chr>, gender <chr>,
## #   race <chr>, court_name <chr>, court_type <chr>, president_name <chr>,
## #   president_party <chr>, nomination_date <chr>, predecessor_last_name <chr>,
## #   predecessor_first_name <chr>, senate_confirmation_date <chr>,
## #   commission_date <chr>, chief_judge_begin <int>, chief_judge_end <int>,
## #   retirement_from_active_service <chr>, termination_date <chr>, ...
```

Data reshaping (`spread()` and `gather()`)

It can be helpful to think of tabular data as coming in two forms: wide data, and long data. Let's load in a table of data. This data contains total membership figures for the Virginia conference of the Methodist Episcopal Church for the years 1812 to 1830.

```
va_wide <- read_csv("http://dh-r.lincolnmullen.com/data/va-methodists-wide.csv")
va_wide
```

```
## # A tibble: 10 x 21
##   conference district `1812` `1813` `1814` `1815` `1816` `1817` `1818` `1819`
##   <chr>      <chr>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Virginia   James Riv~ 5348  4691  4520  4209  4118  3888  3713  3580
## 2 Virginia   Meherren  4882  4486  4771  4687  4702    NA    NA    NA
## 3 Virginia   Meherrin    NA    NA    NA    NA    NA  4435  3964  3860
## 4 Virginia   Neuse      NA    NA  3474  3475  3448  2702  3340  4667
## 5 Virginia   Newbern  3511  3558    NA    NA    NA    NA    NA    NA
```

```
## 6 Virginia Norfolk 4686 6196 6127 6001 5661 6495 6471 NA
## 7 Virginia Raleigh 3822 4018 NA NA NA NA NA NA
## 8 Virginia Roanoke NA NA NA NA 3049 NA 1507 NA
## 9 Virginia Tar River NA NA 3834 3466 NA NA NA NA
## 10 Virginia Yadkin 3174 3216 3528 3323 3374 3323 4689 4547
## # ... with 11 more variables: `1820` <dbl>, `1821` <dbl>, `1822` <dbl>,
## # `1823` <dbl>, `1824` <dbl>, `1825` <dbl>, `1826` <dbl>, `1827` <dbl>,
## # `1828` <dbl>, `1829` <dbl>, `1830` <dbl>
```

The first thing we can notice about this data frame is that it is very wide because it has a column for each of the years. The data is also suitable for reading because it like a table in a publication. We can read from left to right and see when certain districts begin and end and get the values for each year. The difficulties of computing on or plotting the data will also become quickly apparent. How would you make a plot of the change over time in the number of members in each district? Or how would you filter by year, or summarize by year? For that matter, what do the numbers in the table represent, since they are not given an explicit variable name?

The problem with the table is that it is not *tidy data*, because the variables are not in columns and observations in rows. One of the variables is the year, but its values are in the column headers. And another of the variables is total membership, but its values are spread across rows and columns and it is not explicitly named.

The `gather()` function from the `tidyr` package lets us turn wide data into long data. We need to tell the function two kinds of information. First we need to tell it the name of the column to create from the column headers and the name of the implicit variable in the rows. In the example below, we create two new columns `minutes_year` and `total_membership`. Then we also have to tell the function if there are any columns which should remain unchanged. In this case, the `conference` and `district` variables should remain the same, so we remove them from the gathering using the same syntax as the `select()` function.

```
va_wide %>%
  gather(year, members_total, -conference, -district)
```

```
## # A tibble: 190 x 4
##   conference district   year members_total
##   <chr>      <chr>    <chr>         <dbl>
## 1 Virginia James River 1812          5348
## 2 Virginia Meherren 1812          4882
## 3 Virginia Meherrin 1812             NA
## 4 Virginia Neuse 1812             NA
## 5 Virginia Newbern 1812          3511
## 6 Virginia Norfolk 1812          4686
## 7 Virginia Raleigh 1812          3822
## 8 Virginia Roanoke 1812             NA
## 9 Virginia Tar River 1812             NA
## 10 Virginia Yadkin 1812          3174
## # ... with 180 more rows
```

We can see the results above. There are two ways that this result is not quite what we want. Because the years were column headers they are treated as character vectors rather than integers. We can manually convert them in a later step, but we can also let `gather()` do the right thing with the `convert =` argument. Then we have a lot of NA values which were explicit in the wide table but which can be removed from the long table with `na.rm =`.


```
va_long <- va_wide %>%
  gather(year, members_total, -conference, -district,
         convert = TRUE, na.rm = TRUE)
```

```
va_long
```

```
## # A tibble: 100 x 4
##   conference district    year members_total
##   <chr>      <chr>    <int>      <dbl>
## 1 Virginia  James River  1812        5348
## 2 Virginia  Meherren    1812        4882
## 3 Virginia  Newbern     1812        3511
## 4 Virginia  Norfolk     1812        4686
## 5 Virginia  Raleigh     1812        3822
## 6 Virginia  Yadkin      1812        3174
## 7 Virginia  James River 1813        4691
## 8 Virginia  Meherren    1813        4486
## 9 Virginia  Newbern     1813        3558
## 10 Virginia Norfolk      1813        6196
## # ... with 90 more rows
```

Notice that now we can use the data in ggplot2 without any problem.

```
ggplot(va_long,
  aes(x = year, y = members_total, color = district)) +
  geom_line() +
  ggtitle("Membership of districts in the Virginia conference")
```



The inverse operation of `gather()` is `spread()`. With `spread()` we specify the name of the column which should become the new column headers (in this case `minutes_year`), and then the name of the column to fill in underneath those new column headers (in this case, `total_membership`). We can see the results below.

```
va_wide2 <- va_long %>%
  spread(year, members_total)
```

```
va_wide2
```

```
## # A tibble: 10 x 21
##   conference district `1812` `1813` `1814` `1815` `1816` `1817` `1818` `1819`
##   <chr>      <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Virginia James Riv~ 5348 4691 4520 4209 4118 3888 3713 3580
## 2 Virginia Meherren 4882 4486 4771 4687 4702 NA NA NA
## 3 Virginia Meherrin NA NA NA NA NA 4435 3964 3860
## 4 Virginia Neuse NA NA 3474 3475 3448 2702 3340 4667
## 5 Virginia Newbern 3511 3558 NA NA NA NA NA NA
## 6 Virginia Norfolk 4686 6196 6127 6001 5661 6495 6471 NA
## 7 Virginia Raleigh 3822 4018 NA NA NA NA NA NA
## 8 Virginia Roanoke NA NA NA NA 3049 NA 1507 NA
## 9 Virginia Tar River NA NA 3834 3466 NA NA NA NA
## 10 Virginia Yadkin 3174 3216 3528 3323 3374 3323 4689 4547
## # ... with 11 more variables: `1820` <dbl>, `1821` <dbl>, `1822` <dbl>,
## # `1823` <dbl>, `1824` <dbl>, `1825` <dbl>, `1826` <dbl>, `1827` <dbl>,
## # `1828` <dbl>, `1829` <dbl>, `1830` <dbl>
```

By looking at the data we can see that we got back to where we started.

Turning long data into wide is often useful when you want to create a tabular representation of data. (And once you have a data frame that can be a table, the `knitr::kable()` function is quite nice.) And some algorithms, such as clustering algorithms, expect wide data rather than tidy data.

For the exercise, we will use summary statistics of the number of white and black members in the Methodists by year.

```
methodists_by_year_race <- methodists %>%
  group_by(year) %>%
  summarize(white = sum(members_white, na.rm = TRUE),
            black = sum(members_black, na.rm = TRUE),
            indian = sum(members_indian, na.rm = TRUE))
methodists_by_year_race
```

```
## # A tibble: 49 x 4
##   year white black indian
##   <int> <int> <int> <int>
## 1 1786 18291 2890 0
## 2 1787 21949 3883 0
## 3 1788 30557 7991 0
## 4 1789 34425 8840 0
## 5 1790 45983 11682 0
## 6 1791 50580 13098 0
## 7 1792 52079 13871 0
## 8 1793 51486 14420 0
## 9 1794 52794 13906 0
## 10 1795 48121 12171 0
## # ... with 39 more rows
```

- (5) The data in `methodists_by_year_race` could be tidier still. While `white`, `black`, and `indian` are variables, it is perhaps better to think of them as two different variables. One variable would be `race`, containing the racial descriptions that the Methodists used, and another would be `members`, containing

the number of members. Using the `gather()` function, create that data frame.


```
temp <- methodists_by_year_race %>%  
  gather(race, members, -year)
```

temp

```
## # A tibble: 147 x 3  
##   year race  members  
##   <int> <chr>   <int>  
## 1 1786 white  18291  
## 2 1787 white  21949  
## 3 1788 white  30557  
## 4 1789 white  34425  
## 5 1790 white  45983  
## 6 1791 white  50580  
## 7 1792 white  52079  
## 8 1793 white  51486  
## 9 1794 white  52794  
## 10 1795 white  48121  
## # ... with 137 more rows
```

(6) Use the data frame you created in the previous step to create a line plot of membership over time, mapping the `race` column to the `color` aesthetic.

```
ggplot(temp,  
  aes(x = year, y = members, color = race)) +  
  geom_line()
```



- (7) Now use that newly tidied data frame to create a wide data frame, where the years are the column headers and the racial descriptions are the rows.

```
temp2 <- temp %>%
  spread(year, members)
```

```
temp2
```

```
## # A tibble: 3 x 50
##   race   `1786` `1787` `1788` `1789` `1790` `1791` `1792` `1793` `1794` `1795`
##   <chr>   <int> <int> <int> <int> <int> <int> <int> <int> <int> <int>
## 1 black    2890   3883   7991   8840  11682  13098  13871  14420  13906  12171
## 2 indian      0      0      0      0      0      0      0      0      0      0
## 3 white  18291  21949  30557  34425  45983  50580  52079  51486  52794  48121
## # ... with 39 more variables: `1796` <int>, `1797` <int>, `1798` <int>,
## #   `1799` <int>, `1800` <int>, `1801` <int>, `1802` <int>, `1803` <int>,
## #   `1804` <int>, `1805` <int>, `1806` <int>, `1807` <int>, `1808` <int>,
## #   `1809` <int>, `1810` <int>, `1811` <int>, `1812` <int>, `1813` <int>,
## #   `1814` <int>, `1815` <int>, `1816` <int>, `1817` <int>, `1818` <int>,
## #   `1819` <int>, `1820` <int>, `1821` <int>, `1822` <int>, `1823` <int>,
## #   `1824` <int>, `1825` <int>, `1826` <int>, `1827` <int>, `1828` <int>, ...
```

- (8) Now use the same tidied data to create a wide data frame where the racial descriptions are column headers and the years are rows.

```
temp3 <- temp %>%
  spread(race, members)
```

```
temp3
```

```
## # A tibble: 49 x 4
##   year black indian white
##   <int> <int>   <int> <int>
## 1  1786   2890      0 18291
## 2  1787   3883      0 21949
## 3  1788   7991      0 30557
## 4  1789   8840      0 34425
## 5  1790  11682      0 45983
## 6  1791  13098      0 50580
## 7  1792  13871      0 52079
## 8  1793  14420      0 51486
## 9  1794  13906      0 52794
## 10 1795  12171      0 48121
## # ... with 39 more rows
```